

CS7301: Advanced Topics in Optimization for Machine Learning

Lecture 5.2: Coordinate Descent Algorithms

Rishabh Iyer

Department of Computer Science
University of Texas, Dallas

<https://github.com/rishabhk108/AdvancedOptML>

February 26, 2021




Coordinate Descent

-
- Coordinate Descent: Basic Idea
- Optimality w.r.t co-ordinate descent
- Two variants of Coordinate Descent
- Convergence of coordinate descent algorithms.



Coordinate Descent

Goal: Find $\mathbf{x}^* \in \mathbb{R}^d$ minimizing $f(\mathbf{x})$.



Idea: Update one coordinate at a time, while keeping others fixed.

Coordinate Descent

- Modify only one coordinate per step:



Coordinate Descent

- Modify only one coordinate per step:
 - 1 Select $i_k \in [d]$



Coordinate Descent

- Modify only one coordinate per step:

- 1 Select $i_k \in [d]$

- 2 $x_{k+1} = x_k + \gamma \mathbf{e}_{i_k}$



Coordinate Descent

- Modify only one coordinate per step:
 - ① Select $i_k \in [d]$
 - ② $x_{k+1} = x_k + \gamma e_{i_k}$
- γ is the step size of coordinate descent



Coordinate Descent

- Modify only one coordinate per step:
 - 1 Select $i_k \in [d]$
 - 2 $x_{k+1} = x_k + \gamma e_{i_k}$
- γ is the step size of coordinate descent
- Two variants of coordinate descent:



Coordinate Descent

- Modify only one coordinate per step:
 - ① Select $i_k \in [d]$
 - ② $x_{k+1} = x_k + \gamma e_{i_k}$
- γ is the step size of coordinate descent
- Two variants of coordinate descent:
- Gradient based step size:

$$x_{k+1} = x_k - \frac{1}{L} \nabla_{i_k} f(x_k) e_{i_k}$$



Coordinate Descent

- Modify only one coordinate per step:
 - 1 Select $i_k \in [d]$
 - 2 $x_{k+1} = x_k + \gamma e_{i_k}$
- γ is the step size of coordinate descent
- Two variants of coordinate descent:
- Gradient based step size:

$$x_{k+1} = x_k - \frac{1}{L} \nabla_{i_k} f(x_k) e_{i_k}$$

- Exact coordinate minimization: solve the single variable minimization

$$\operatorname{argmin}_{\gamma \in \mathbb{R}} f(x_k + \gamma e_{i_k})$$

in closed form.



Coordinate Descent

- Two variants of coordinate descent:



Coordinate Descent

- Two variants of coordinate descent:
- Gradient based step size:

$$x_{k+1} = x_k - \frac{1}{L} \nabla_{i_k} f(x_k) e_{i_k}$$



Coordinate Descent

- Two variants of coordinate descent:
- Gradient based step size:

$$x_{k+1} = x_k - \frac{1}{L} \nabla_{i_k} f(x_k) e_{i_k}$$

- Exact coordinate minimization: solve the single variable minimization

$$\operatorname{argmin}_{\gamma \in \mathbb{R}} f(x_k + \gamma e_{i_k})$$

in closed form.



Coordinate Descent

- Two variants of coordinate descent:
- Gradient based step size:

$$x_{k+1} = x_k - \frac{1}{L} \nabla_{i_k} f(x_k) e_{i_k}$$

- Exact coordinate minimization: solve the single variable minimization

$$\operatorname{argmin}_{\gamma \in \mathbb{R}} f(x_k + \gamma e_{i_k})$$

in closed form.

- Also how do we select the coordinate i_k ?



Coordinate Descent

- Two variants of coordinate descent:
- Gradient based step size:

$$x_{k+1} = x_k - \frac{1}{L} \nabla_{i_k} f(x_k) e_{i_k}$$

- Exact coordinate minimization: solve the single variable minimization

$$\operatorname{argmin}_{\gamma \in \mathbb{R}} f(x_k + \gamma e_{i_k})$$

in closed form.

- Also how do we select the coordinate i_k ?
- Two strategies: One is to pick it up randomly (called randomized coordinate descent) and second it to pick it up greedily (greedy coordinate descent).



Randomized Coordinate Descent

select $i_t \in [d]$ uniformly at random
 $\mathbf{x}_{t+1} := \mathbf{x}_t - \frac{1}{L} \nabla_{i_t} f(\mathbf{x}_t) \mathbf{e}_{i_t}$

- **Faster convergence** than gradient descent
(if coordinate step is significantly cheaper than full gradient step)



Randomized Coordinate Descent: Convergence Analysis

- Key ingredient: **Coordinate wise Smoothness**:

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$



Randomized Coordinate Descent: Convergence Analysis

- Key ingredient: **Coordinate wise Smoothness**:

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$

- This is equivalent to coordinate wise Lipschitz gradient:

$$|\nabla_i f(x + \gamma e_i) - \nabla_i f(x)| \leq L_i |\gamma|$$



Randomized Coordinate Descent: Convergence Analysis

- Key ingredient: **Coordinate wise Smoothness**:

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$

- This is equivalent to coordinate wise Lipschitz gradient:

$$|\nabla_i f(x + \gamma e_i) - \nabla_i f(x)| \leq L_i |\gamma|$$

- Define $L = \max_i L_i$ as the Maximum coordinate wise Lipschitz smoothness



Randomized Coordinate Descent: Convergence Analysis

- Key ingredient: **Coordinate wise Smoothness**:

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$

- This is equivalent to coordinate wise Lipschitz gradient:

$$|\nabla_i f(x + \gamma e_i) - \nabla_i f(x)| \leq L_i |\gamma|$$

- Define $L = \max_i L_i$ as the Maximum coordinate wise Lipschitz smoothness
- Similar to previous analysis, let's assume $\|x_0 - x^*\|^2 \leq R^2$.



Randomized Coordinate Descent: Convergence Result

- Key ingredient: **Coordinate wise Smoothness**:

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$



Randomized Coordinate Descent: Convergence Result

- Key ingredient: **Coordinate wise Smoothness**:

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$

- Define $L = \max_i L_i$ as the Maximum coordinate wise Lipschitz smoothness. Also, assume $\|x_0 - x^*\|^2 \leq R^2$.



Randomized Coordinate Descent: Convergence Result

- Key ingredient: **Coordinate wise Smoothness:**

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$

- Define $L = \max_i L_i$ as the Maximum coordinate wise Lipschitz smoothness. Also, assume $\|x_0 - x^*\|^2 \leq R^2$.
- **Theorem:** For a coordinate wise smooth convex function, we have:

$$E(f(x_k)) - f_* \leq \frac{2dLR_0^2}{k}$$



Randomized Coordinate Descent: Convergence Result

- Key ingredient: **Coordinate wise Smoothness:**

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$

- Define $L = \max_i L_i$ as the Maximum coordinate wise Lipschitz smoothness. Also, assume $\|x_0 - x^*\|^2 \leq R^2$.
- **Theorem:** For a coordinate wise smooth convex function, we have:

$$E(f(x_k)) - f_* \leq \frac{2dLR_0^2}{k}$$

- Similar to Gradient Descent, this is $O(1/\epsilon)$ convergence!



Randomized Coordinate Descent: Convergence Result

- Key ingredient: **Coordinate wise Smoothness:**

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$

- Define $L = \max_i L_i$ as the Maximum coordinate wise Lipschitz smoothness. Also, assume $\|x_0 - x^*\|^2 \leq R^2$.
- **Theorem:** For a coordinate wise smooth convex function, we have:

$$E(f(x_k)) - f_* \leq \frac{2dLR_0^2}{k}$$

- Similar to Gradient Descent, this is $O(1/\epsilon)$ convergence!
- However, we just need to update a single dimension at a time, and hence the per iteration cost of CD can be $O(d)$ cheaper compared to GD!



Randomized Coordinate Descent: Convergence Result

- Key ingredient: **Coordinate wise Smoothness**:

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$



Randomized Coordinate Descent: Convergence Result

- Key ingredient: **Coordinate wise Smoothness**:

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$

- Define $L = \max_i L_i$ as the Maximum coordinate wise Lipschitz smoothness. Also, assume $\|x_0 - x^*\|^2 \leq R^2$.



Randomized Coordinate Descent: Convergence Result

- Key ingredient: **Coordinate wise Smoothness**:

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$

- Define $L = \max_i L_i$ as the Maximum coordinate wise Lipschitz smoothness. Also, assume $\|x_0 - x^*\|^2 \leq R^2$.
- If additionally, f is μ -strongly convex? (Recall strong convexity implies $f(y) \geq f(x) + \nabla f(x)^T (y - x) + \frac{\mu}{2} \|y - x\|^2$, for all x, y)



Randomized Coordinate Descent: Convergence Result

- Key ingredient: **Coordinate wise Smoothness**:

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$

- Define $L = \max_i L_i$ as the Maximum coordinate wise Lipschitz smoothness. Also, assume $\|x_0 - x^*\|^2 \leq R^2$.
- If additionally, f is μ -strongly convex? (Recall strong convexity implies $f(y) \geq f(x) + \nabla f(x)^T (y - x) + \frac{\mu}{2} \|y - x\|^2$, for all x, y)
- **Theorem:** If f is coordinate wise Lipschitz Smooth + Strongly Convex, we have:

$$E(f(x_k)) - f_* \leq \left(1 - \frac{\mu}{dL}\right)^k (f(x_0) - f_*)$$



Randomized Coordinate Descent: Improvement at Every iteration!

- Key ingredient: **Coordinate wise Smoothness:**

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$



Randomized Coordinate Descent: Improvement at Every iteration!

- Key ingredient: **Coordinate wise Smoothness**:

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$

- Define $L = \max_i L_i$ as the Maximum coordinate wise Lipschitz smoothness. Also, assume $\|x_0 - x^*\|^2 \leq R^2$.



Randomized Coordinate Descent: Improvement at Every iteration!

- Key ingredient: **Coordinate wise Smoothness**:

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$

- Define $L = \max_i L_i$ as the Maximum coordinate wise Lipschitz smoothness. Also, assume $\|x_0 - x^*\|^2 \leq R^2$.
- It holds that:

$$f(x_{k+1}) \leq f(x_k) - \frac{1}{2L} [\nabla f(x_k)]_{i_k}^2$$



Randomized Coordinate Descent: Improvement at Every iteration!

- Key ingredient: **Coordinate wise Smoothness:**

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$

- Define $L = \max_i L_i$ as the Maximum coordinate wise Lipschitz smoothness. Also, assume $\|x_0 - x^*\|^2 \leq R^2$.
- It holds that:

$$f(x_{k+1}) \leq f(x_k) - \frac{1}{2L} [\nabla f(x_k)]_{i_k}^2$$

- This means that similar to GD, CD improves the objective value at every iteration!



Importance Sampling

- Key ingredient: **Coordinate wise Smoothness**:

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$



Importance Sampling

- Key ingredient: **Coordinate wise Smoothness**:

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$

- Instead of random selection, can we be slightly smarter?



Importance Sampling

- Key ingredient: **Coordinate wise Smoothness**:

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$

- Instead of random selection, can we be slightly smarter?
- Sample i_k with probability: $P[i_k = i] = \frac{L_i}{\sum_i L_i}$, and use step size $1/L_{i_k}$, we can get a slightly better rate of convergence



Importance Sampling

- Key ingredient: **Coordinate wise Smoothness**:

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$

- Instead of random selection, can we be slightly smarter?
- Sample i_k with probability: $P[i_k = i] = \frac{L_i}{\sum_i L_i}$, and use step size $1/L_{i_k}$, we can get a slightly better rate of convergence
 - ① Smooth Functions: $\frac{2d\bar{L}R_0^2}{k}$



Importance Sampling

- Key ingredient: **Coordinate wise Smoothness**:

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$

- Instead of random selection, can we be slightly smarter?
- Sample i_k with probability: $P[i_k = i] = \frac{L_i}{\sum_i L_i}$, and use step size $1/L_{i_k}$, we can get a slightly better rate of convergence
 - 1 Smooth Functions: $\frac{2d\bar{L}R_0^2}{k}$
 - 2 Smooth + Strongly Convex: $(1 - \frac{\mu}{dL})^k (f(x_0) - f_*)$



Importance Sampling

- Key ingredient: **Coordinate wise Smoothness**:

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$

- Instead of random selection, can we be slightly smarter?
- Sample i_k with probability: $P[i_k = i] = \frac{L_i}{\sum_i L_i}$, and use step size $1/L_{i_k}$, we can get a slightly better rate of convergence
 - ① Smooth Functions: $\frac{2d\bar{L}R_0^2}{k}$
 - ② Smooth + Strongly Convex: $(1 - \frac{\mu}{dL})^k (f(x_0) - f_*)$
- In the above, $\bar{L} = \sum_i L_i / d$



Importance Sampling

- Key ingredient: **Coordinate wise Smoothness**:

$$f(x + \gamma e_i) \leq f(x) + \gamma \nabla_i f(x) + \frac{L_i}{2} \gamma^2, \forall x \in \mathbb{R}^d, \forall \gamma \in \mathbb{R}$$

- Instead of random selection, can we be slightly smarter?
- Sample i_k with probability: $P[i_k = i] = \frac{L_i}{\sum_i L_i}$, and use step size $1/L_{i_k}$, we can get a slightly better rate of convergence
 - 1 Smooth Functions: $\frac{2d\bar{L}R_0^2}{k}$
 - 2 Smooth + Strongly Convex: $(1 - \frac{\mu}{dL})^k (f(x_0) - f_*)$
- In the above, $\bar{L} = \sum_i L_i / d$
- Note that $\bar{L} \leq L$ and can in fact be significantly smaller!



Steepest (or Greedy) Coordinate Descent

- Instead of random or importance sampling based selection, select greedily!



Steepest (or Greedy) Coordinate Descent

- Instead of random or importance sampling based selection, select greedily!
- Select:

$$i_k = \operatorname{argmax}_{i \in [d]} [\nabla_i f(x_k)]$$



Steepest (or Greedy) Coordinate Descent

- Instead of random or importance sampling based selection, select greedily!
- Select:

$$i_k = \operatorname{argmax}_{i \in [d]} [\nabla_i f(x_k)]$$

- Note, this strategy is deterministic instead of random.



Steepest (or Greedy) Coordinate Descent

- Instead of random or importance sampling based selection, select greedily!
- Select:

$$i_k = \operatorname{argmax}_{i \in [d]} [\nabla_i f(x_k)]$$

- Note, this strategy is deterministic instead of random.
- Vanilla analysis gives the same worst case Convergence rate as the randomized case.



Steepest (or Greedy) Coordinate Descent

- Instead of random or importance sampling based selection, select greedily!

- Select:

$$i_k = \operatorname{argmax}_{i \in [d]} [\nabla_i f(x_k)]$$

- Note, this strategy is deterministic instead of random.
- Vanilla analysis gives the same worst case Convergence rate as the randomized case.
- However, if we assume f is strongly convex w.r.t the l_1 norm with parameter $\mu_1 > 0$, we get improved rates of:



Steepest (or Greedy) Coordinate Descent

- Instead of random or importance sampling based selection, select greedily!
- Select:

$$i_k = \operatorname{argmax}_{i \in [d]} [\nabla_i f(x_k)]$$

- Note, this strategy is deterministic instead of random.
- Vanilla analysis gives the same worst case Convergence rate as the randomized case.
- However, if we assume f is strongly convex w.r.t the l_1 norm with parameter $\mu_1 > 0$, we get improved rates of:
 - ① Smooth Functions: $\frac{2LR_0^2}{k}$



Steepest (or Greedy) Coordinate Descent

- Instead of random or importance sampling based selection, select greedily!
- Select:

$$i_k = \operatorname{argmax}_{i \in [d]} [\nabla_i f(x_k)]$$

- Note, this strategy is deterministic instead of random.
- Vanilla analysis gives the same worst case Convergence rate as the randomized case.
- However, if we assume f is strongly convex w.r.t the l_1 norm with parameter $\mu_1 > 0$, we get improved rates of:
 - 1 Smooth Functions: $\frac{2LR_0^2}{k}$
 - 2 Smooth + Strongly Convex: $(1 - \frac{\sigma}{L})^k (f(x_0) - f_*)$



Exact Coordinate Minimization

- So far, we went over coordinate descent in a way akin to gradient descent, i.e. taking a coordinate step in each coordinate.
- Next, consider exact minimization in each coordinate.

$$x_1^{(k)} \in \operatorname{argmin}_{x_1} f(x_1, x_2^{(k-1)}, x_3^{(k-1)}, \dots, x_n^{(k-1)})$$

$$x_2^{(k)} \in \operatorname{argmin}_{x_2} f(x_1^{(k)}, x_2, x_3^{(k-1)}, \dots, x_n^{(k-1)})$$

$$x_3^{(k)} \in \operatorname{argmin}_{x_3} f(x_1^{(k)}, x_2^{(k)}, x_3, \dots, x_n^{(k-1)})$$

...

$$x_n^{(k)} \in \operatorname{argmin}_{x_n} f(x_1^{(k)}, x_2^{(k)}, x_3^{(k)}, \dots, x_n)$$

for $k = 1, 2, 3, \dots$

Note: after we solve for $x_i^{(k)}$, we use its new value from then on!

Exact Coordinate Minimization: Linear Regression

Consider linear regression

$$\min_{\beta \in \mathbb{R}^p} \frac{1}{2} \|y - X\beta\|_2^2$$

where $y \in \mathbb{R}^n$, and $X \in \mathbb{R}^{n \times p}$ with columns X_1, \dots, X_p

Minimizing over β_i , with all β_j , $j \neq i$ fixed:

$$0 = \nabla_i f(\beta) = X_i^T (X\beta - y) = X_i^T (X_i\beta_i + X_{-i}\beta_{-i} - y)$$

i.e., we take


$$\beta_i = \frac{X_i^T (y - X_{-i}\beta_{-i})}{X_i^T X_i}$$

Coordinate descent repeats this update for $i = 1, 2, \dots, p, 1, 2, \dots$

LAS

Exact Coordinate Minimization: Linear Regression

Coordinate descent vs gradient descent for linear regression: 100 instances ($n = 100$, $p = 20$)



Is it fair to compare 1 cycle of coordinate descent to 1 iteration of gradient descent? Yes, if we're clever:

$$\beta_i \leftarrow \frac{X_i^T (y - X_{-i} \beta_{-i})}{X_i^T X_i} = \frac{X_i^T r}{\|X_i\|_2^2} + \beta_i$$

where $r = y - X\beta$. Therefore each coordinate update takes $O(n)$ operations — $O(n)$ to update r , and $O(n)$ to compute $X_i^T r$ — and one cycle requires $O(np)$ operations, just like gradient descent


DALLAS

Optimality Conditions with Coordinate Descent

- Coordinate Descent Family of Algorithms try to make progress every iteration!
- What is the stopping/optimality condition?
- One way to think of this is: If we are a point x such that it is minimum along each coordinate axis! In other words, no coordinate descent algorithm can make progress!
- Does this mean we are at a global minimum?
- Mathematically this means: Does a point x satisfying $f(x + \delta e_i) \geq f(x), \forall \delta, i \Rightarrow f(x) = \min_z f(z)$?



Optimality Conditions: Differentiable Functions



A: Yes! Proof:



$$\nabla f(x) = \left(\frac{\partial f}{\partial x_1}(x), \dots, \frac{\partial f}{\partial x_n}(x) \right) = 0$$

Q: Same question, but for f convex (not differentiable) ... ?

LAS

17/20

Optimality Conditions: Non Differentiable Functions





A: No! Look at the above counterexample

Q: Same question again, but now $f(x) = g(x) + \sum_{i=1}^n h_i(x_i)$, with g convex, differentiable and each h_i convex ... ? (Nonsmooth part here called **separable**)

LAS

Optimality Conditions: Seperable Functions



A: Yes! Proof: for any y ,

$$\begin{aligned} f(y) - f(x) &\geq \nabla g(x)^T (y - x) + \sum_{i=1}^n [h_i(y_i) - h_i(x_i)] \\ &= \sum_{i=1}^n \underbrace{[\nabla_i g(x)(y_i - x_i) + h_i(y_i) - h_i(x_i)]}_{\geq 0} \geq 0 \end{aligned}$$

LAS

Coordinate Descent Summary

- Algorithms:



Coordinate Descent Summary

- Algorithms:
 - 1 Randomized Coordinate Descent



Coordinate Descent Summary

- Algorithms:
 - ① Randomized Coordinate Descent
 - ② Greedy (Steepest) Coordinate Descent



Coordinate Descent Summary

- Algorithms:
 - ① Randomized Coordinate Descent
 - ② Greedy (Steepest) Coordinate Descent
 - ③ Exact Coordinate Minimization



Coordinate Descent Summary

- Algorithms:
 - ① Randomized Coordinate Descent
 - ② Greedy (Steepest) Coordinate Descent
 - ③ Exact Coordinate Minimization
- Convergence Results:



Coordinate Descent Summary

- Algorithms:

- ① Randomized Coordinate Descent
- ② Greedy (Steepest) Coordinate Descent
- ③ Exact Coordinate Minimization

- Convergence Results:

- ① Randomized Coordinate Descent: Smooth Functions ($\frac{2dLR_0^2}{k}$)



Coordinate Descent Summary

- Algorithms:

- 1 Randomized Coordinate Descent
- 2 Greedy (Steepest) Coordinate Descent
- 3 Exact Coordinate Minimization

- Convergence Results:

- 1 Randomized Coordinate Descent: Smooth Functions ($\frac{2dLR_0^2}{k}$)
- 2 Smooth + Strongly Convex: $(1 - \frac{\sigma}{dL})^k (f(x_0) - f_*)$



Coordinate Descent Summary

- Algorithms:

- ① Randomized Coordinate Descent
- ② Greedy (Steepest) Coordinate Descent
- ③ Exact Coordinate Minimization

- Convergence Results:

- ① Randomized Coordinate Descent: Smooth Functions ($\frac{2dLR_0^2}{k}$)
- ② Smooth + Strongly Convex: $(1 - \frac{\sigma}{dL})^k (f(x_0) - f_*)$
- ③ Slight Improvement for Steepest CD: $(1 - \frac{\sigma}{dL})^k (f(x_0) - f_*)$



Coordinate Descent Summary

- Algorithms:

- 1 Randomized Coordinate Descent
- 2 Greedy (Steepest) Coordinate Descent
- 3 Exact Coordinate Minimization

- Convergence Results:

- 1 Randomized Coordinate Descent: Smooth Functions ($\frac{2dLR_0^2}{k}$)
- 2 Smooth + Strongly Convex: $(1 - \frac{\sigma}{dL})^k (f(x_0) - f_*)$
- 3 Slight Improvement for Steepest CD: $(1 - \frac{\sigma}{dL})^k (f(x_0) - f_*)$
- 4 Similar bound for exact coordinate minimization. Method of choice if the coordinate wise minimization can be solved exactly!



Coordinate Descent Summary

- Algorithms:

- 1 Randomized Coordinate Descent
- 2 Greedy (Steepest) Coordinate Descent
- 3 Exact Coordinate Minimization

- Convergence Results:

- 1 Randomized Coordinate Descent: Smooth Functions ($\frac{2dLR_0^2}{k}$)
- 2 Smooth + Strongly Convex: $(1 - \frac{\sigma}{dL})^k (f(x_0) - f_*)$
- 3 Slight Improvement for Steepest CD: $(1 - \frac{\sigma}{dL})^k (f(x_0) - f_*)$
- 4 Similar bound for exact coordinate minimization. Method of choice if the coordinate wise minimization can be solved exactly!
- 5 Above bounds can also be extended if the non-differentiability is separable (e.g. L1 regularization).



Coordinate Descent Summary

- Algorithms:

- 1 Randomized Coordinate Descent
- 2 Greedy (Steepest) Coordinate Descent
- 3 Exact Coordinate Minimization

- Convergence Results:

- 1 Randomized Coordinate Descent: Smooth Functions ($\frac{2dLR_0^2}{k}$)
- 2 Smooth + Strongly Convex: $(1 - \frac{\sigma}{dL})^k (f(x_0) - f_*)$
- 3 Slight Improvement for Steepest CD: $(1 - \frac{\sigma}{dL})^k (f(x_0) - f_*)$
- 4 Similar bound for exact coordinate minimization. Method of choice if the coordinate wise minimization can be solved exactly!
- 5 Above bounds can also be extended if the non-differentiability is separable (e.g. L1 regularization).

- Extensions: Accelerated Coordinate Minimization also possible!

